

Credit Card Default Prediction

Machine learning approach

Adarsh Yadav - 22112007



Problem Overview

Banks are facing problem with credit card default so they want to know whether the customer will default next month or not



Solution Approach

We have the data of almost 25000 customers so based on this data we will train our machine learning classification model .



Dataset Overview

Summary:

Dataset of **25,000 credit card users** across 6 months, with attributes like payment history, demographic data, and default status.

Key Columns:

Column Name	Description
CustomerId	Unique ID per customer
Marriage	Marital Status: 1=Married, 2=Single, 3=Others
Sex	Gender: 1=Male, 0=Female
Education	1=Graduate School, 2=University, 3=High School, 4=Others
Limit_balance	Credit limit (in currency units)
Age	Age of customer



Dataset Overview

Monthly Behavior:

- **PAY_X (X = 0 to 6):** Payment status
 - **-2**: No credit use
 - **-1**: Full repayment
 - **0**: Minimum/partial payment
 - **1+**: Delayed payments
- **BILL_AMT_X (X = 1 to 6):** Bill amount
 - **> 0**: Money owed
 - **= 0**: No activity
 - **< 0**: Overpaid (credit balance)
- **PAY_AMT_X (X = 1 to 6):** Payments made

Engineered Features:

Column Name	Description
Avg_bill_amt	Average bill across 6 months
PAY_TO_BILL_ratio	Total payment / total bill (6 months)

Target:

Column	Description
next_month_default	1 : Will default next month 0 : Will not default

```
✓ 0s [▶] df.shape  
[↔] (25121, 27)
```



Dataset Overview

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25247 entries, 0 to 25246
Data columns (total 27 columns):

#	Column	Non-Null Count	Dtype
0	Customer_ID	25247 non-null	int64
1	marriage	25247 non-null	int64
2	sex	25247 non-null	int64
3	education	25247 non-null	int64
4	LIMIT_BAL	25247 non-null	int64
5	age	25121 non-null	float64
6	pay_0	25247 non-null	int64
7	pay_2	25247 non-null	int64
8	pay_3	25247 non-null	int64
9	pay_4	25247 non-null	int64
10	pay_5	25247 non-null	int64
11	pay_6	25247 non-null	int64
12	Bill_amt1	25247 non-null	float64
13	Bill_amt2	25247 non-null	float64
14	Bill_amt3	25247 non-null	float64
15	Bill_amt4	25247 non-null	float64
16	Bill_amt5	25247 non-null	float64
17	Bill_amt6	25247 non-null	float64
18	pay_amt1	25247 non-null	float64
19	pay_amt2	25247 non-null	float64
20	pay_amt3	25247 non-null	float64
21	pay_amt4	25247 non-null	float64
22	pay_amt5	25247 non-null	float64
23	pay_amt6	25247 non-null	float64
24	AVG_Bill_amt	25247 non-null	float64
25	PAY_TO_BILL_ratio	25247 non-null	float64
26	next_month_default	25247 non-null	int64

dtypes: float64(15), int64(12)
memory usage: 5.2 MB



Exploratory Data Analysis

Age: There are some null values around 126 so removing those rows as they less so it will not affect the result

Removing rows with null values

```
df.isnull().sum()  
# Age is having some null values
```

	0
Customer_ID	0
marriage	0
sex	0
education	0
LIMIT_BAL	0
age	126

```
[34] # removing them  
df.dropna(subset=['age'],inplace=True)
```




Replacing Categorical Columns with original value for better understanding

```
df['sex'] = df['sex'].replace({1:'Male',0:'Female'})  
df['education'] = df['education'].replace({1:'Graduate School', 2:'University', 3:'High School',0:'Others',4:'Others',5:'Others',6:'Others' })  
df['marriage'] = df['marriage'].replace({1:'Married',2:'No',3:'Others',0:'Others'})  
df.head()
```



Age :-




Observations:

- Age is between 21 to 79
- Mostly are between 25 to 35
- Age data is rightly skewed

```
# Skewness  
df['age'].skew()  
  
np.float64(0.7384976486065989)
```

Age :- Multivariate Analysis

Age with Limit Balance





Observations:

- Shows that if age and limit balance are higher chance of defaulting is low



Age :- Multivariate Analysis

Age with Pay amount

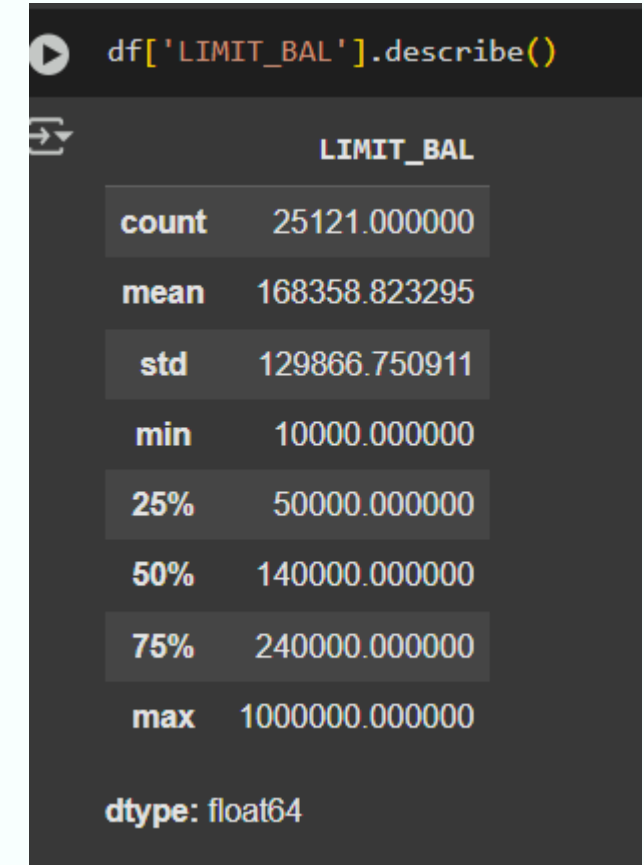




Observations:

- Customer with high payment amount are mostly between 25-50 ,
- Below 35 - Defaulter are decreasing as we move from month1 to month6



Limit Balance :-






Observations:

- Most Customer have limit balance between 0-175000
- Default rate is higher below 175000
- Rightly skewed



Limit Balance :-Multivariate

Limit Balance with Sex



Observations:

- Mean Limit is almost same for both gender
- When separate on the basis of default, mean limits differs significantly for defaulters and non defaulters

```
df.groupby(['sex'])['LIMIT_BAL'].describe(percentiles=[.1,.2,.3,.4,.5,.6,.7,.8,.9,.99])
```



	count	mean	std	min	10%	20%	30%	40%	50%	60%	70%	80%	90%	99%	max
sex															
Female	9930.0	164483.987915	136518.286753	10000.0	20000.0	50000.0	50000.0	90000.0	130000.0	170000.0	210000.0	280000.0	360000.0	500000.0	800000.0
Male	15191.0	170891.712198	125267.964752	10000.0	30000.0	50000.0	80000.0	110000.0	150000.0	180000.0	220000.0	270000.0	360000.0	500000.0	1000000.0

```
[52] df.groupby(['sex','next_month_default'])['LIMIT_BAL'].describe(percentiles=[.1,.2,.3,.4,.5,.6,.7,.8,.9,.99])
```

		count	mean	std	min	10%	20%	30%	40%	50%	60%	70%	80%	90%	99%
sex	next_month_default														
Female	0	7858.0	174846.780351	139523.143559	10000.0	30000.0	50000.0	60000.0	100000.0	150000.0	180000.0	230000.0	290000.0	380000.0	500000.0
	1	2072.0	125183.397683	116390.114564	10000.0	20000.0	30000.0	50000.0	50000.0	80000.0	120000.0	160000.0	200000.0	300000.0	500000.0
Male	0	12479.0	179229.585704	126337.685866	10000.0	40000.0	60000.0	90000.0	120000.0	160000.0	200000.0	230000.0	280000.0	360000.0	500000.0
	1	2712.0	132525.811209	112545.737566	10000.0	30000.0	40000.0	50000.0	70000.0	90000.0	130000.0	170000.0	220000.0	300000.0	500000.0

Limit Balance :-

Limit Balance with Bill amount





Observations:

- Number of customer drastically decreases after 500000
- Limit balance increase and bill amount decreases



Bill Amount:-






Observations:

- Rightly skewed data
- There is some relation from month 1 to month 6



Bill Amount:-Multivariate

Bill amount with sex and default



Observations:

- From the graphs we might come to an assumption that bill amounts has a significance effect on gender.
- Mean amouns between male and female shows some difference

```
for col in bill_amt_cols:
    print(col )
    print(df.groupby(['sex'])[col].describe(percentiles=[.1,.2,.3,.4,.5,.6,.7,.8,.9,.99]))
```

Bill_amt1	count	mean	std	min	10%	20%	30%	\
sex								
Female	9930.0	54177.635872	76854.252962	0.0	389.798	2501.386	7722.329	
Male	15191.0	49266.369535	70878.224942	0.0	139.860	1603.630	5024.350	

Bill Amount:-Multivariate

Bill amount with education

```
for col in bill_amt_cols:
    print(col)
    print(df.groupby(['education'])[col].describe(percentiles=[.1,.2,.3,.4,.5,.6,.7,.8,.9,.99]))
```

Bill_amt1	count	mean	std	min	10%	20%	\
education							
Graduate School	8944.0	48717.536096	78350.806939	0.0	0.590	795.058	
High School	4096.0	47455.568801	64983.865689	0.0	389.730	2565.270	
Others	424.0	70014.542524	88573.547889	0.0	319.795	2955.308	
University	11657.0	53752.725842	71297.852296	0.0	491.218	3533.254	

	30%	40%	50%	60%	70%	\
education						
Graduate School	2773.729	6457.104	14217.43	26854.704	47926.980	
High School	7824.655	15968.880	24930.56	38243.410	50392.780	
Others	8351.540	17872.266	30639.68	53840.382	88624.442	
University	9891.824	18156.004	27527.00	44303.040	58206.748	



	80%	90%	99%	max
education				
Graduate School	81251.568	149612.218	376876.4477	964511.16
High School	75268.330	126646.550	285091.1735	746813.18
Others	136150.440	192473.288	372203.4185	626647.12
University	86710.552	140750.082	339377.2104	610722.35

Observations:

- Median bill amounts of each educational category, displays only a significantly small difference as it passes through each month
- When each category is subcategorized based on 'default' value, we see that the median bill amount for 'Others' category shows some difference for defaulters and non defaulters in every month.



Sex:-



Observations:

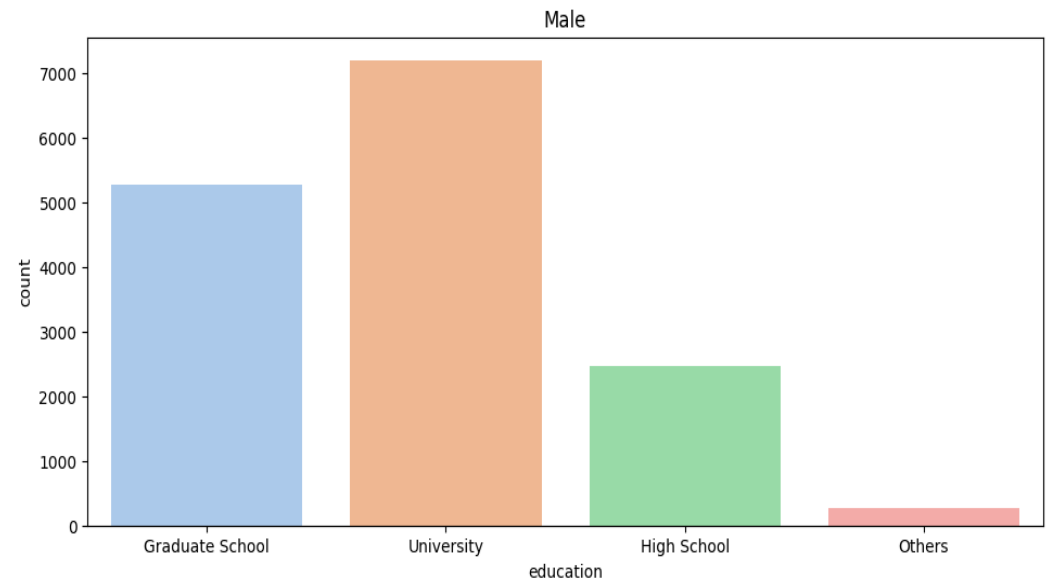
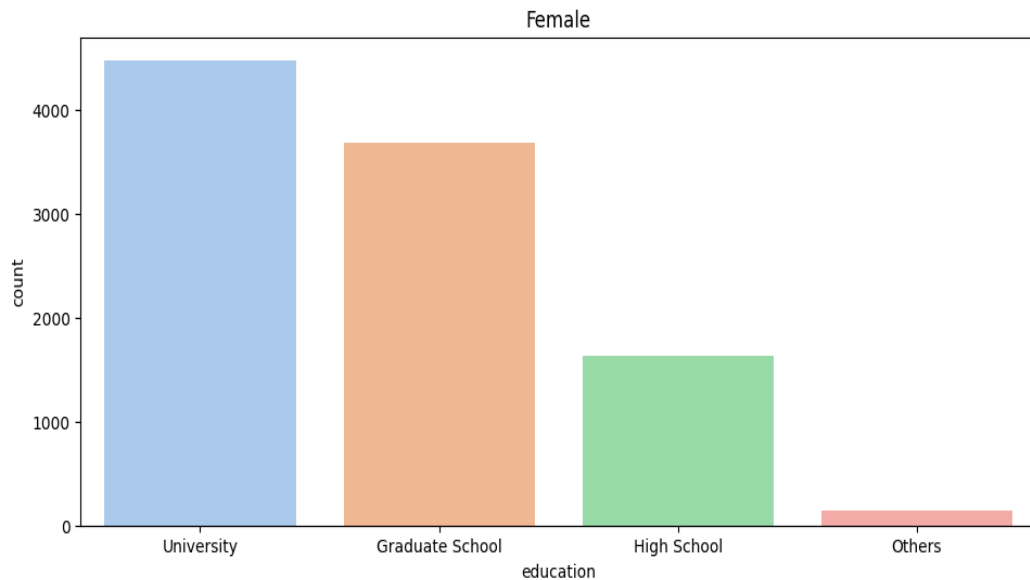
- More Male customers than female
- More percentage of female defaulters than male.

```
[61] df['sex'].value_counts()
```

count	
sex	
Male	15191
Female	9930

Sex:- Multivariate

Sex with Education



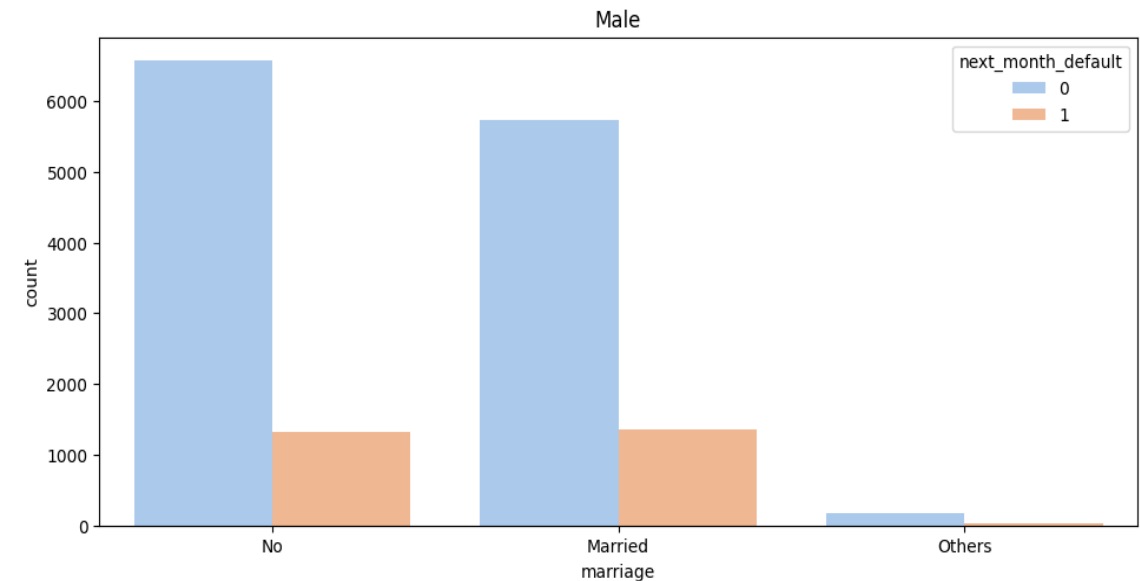

Observations:

- Among Both gender mostly are university educated



Sex:- Multivariate

Sex with Marriage





Observations:

- Both married and unmarried almost have same amount of defaulters



Sex:- Multivariate

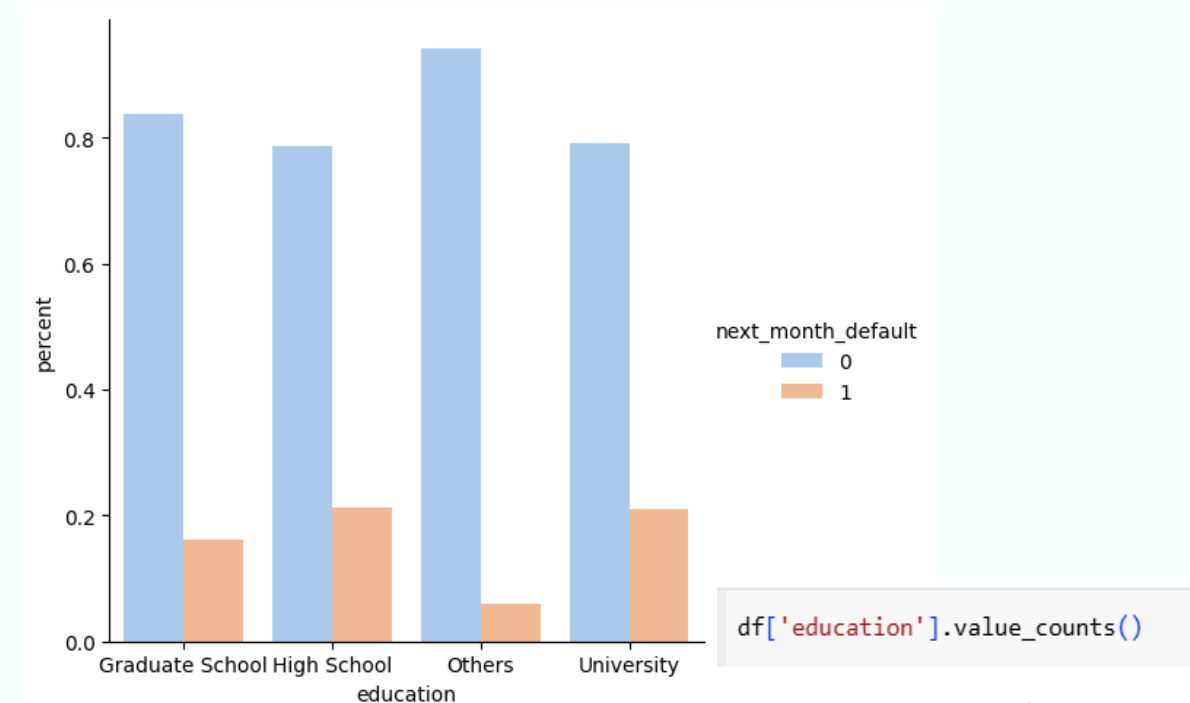

Sex with Repayment



Observations:

- the number of defaulters have seems to be almost on same level for the first 5 months, but during the last month has shown a small dip.

Education:-Univariate



```
df['education'].value_counts()
```

Observations:



- University students are the majority of customers
- Majority of defaulter are from University and High School

count	
education	
University	11657
Graduate School	8944
High School	4096
Others	424



Education:-Multivariate

Education with Marriage




Observations:

- Proportion of defaulter to Non defaulter is almost same in all education categories



Education:-Multivariate

Education with Repayment status




Observations:

- We can see that, for every education category when we move from Month1 to Month6, there seems to be a sudden increase in the number of customer with 1 month repayment delay in the month 6, when all the other months, there were almost none.




Marriage:-Univariate



Observations:

- Unmarried people are more in count
- Default proportion is almost same in all categories



```
df['marriage'].value_counts()
```



count

marriage

No 13374


Married 11424

Others 323



Marriage:-Multivariate

Marriage with repayment amount



Observations:

- Trend is almost same above categories with small variation

Feature Engineering:-

Removing Null values

✓ Removing Null value

```
[3] ### Dropping age with null values as they are only 126 and dropping them will not effect data  
df.dropna(subset=['age'],inplace=True)
```



Feature Engineering:-

Handling Class Imbalance using SMOTE-

Synthetic Minority Oversampling

Technique is a machine learning technique used to address the problem of imbalanced datasets. It works by **creating synthetic data** points for the minority class, effectively increasing its representation in the dataset and helping models learn to classify it more accurately.

```
[4] smote = SMOTE()
```

```
[5] X = df.iloc[:, :-1]  
    y = df['next_month_default']
```

```
[6] x_smote, y_smote = smote.fit_resample(X, y)
```

```
[7] df_final = pd.DataFrame(x_smote, columns=df.columns[:-1])  
    df_final['next_month_default'] = y_smote  
  
    df_final.head()
```

Feature Engineering:-

Changing Column Names :

✓ Replacing Months with original name

[+ Code](#)[+ Text](#)

```
[13] df_final.rename(columns= {'pay_6':'Pay_January','pay_5':'Pay_February','pay_4':'Pay_March','pay_3':'Pay_April','pay_2':'Pay_May','pay_1':'Pay_June'})
df_final.rename(columns= {'pay_amt6':'Pay_amt_January','pay_amt5':'Pay_amt_February','pay_amt4':'Pay_amt_March','pay_amt3':'Pay_amt_April','pay_amt2':'Pay_amt_May','pay_amt1':'Pay_amt_June'})
df_final.rename(columns= {'Bill_amt6':'Bill_amt_January','Bill_amt5':'Bill_amt_February','Bill_amt4':'Bill_amt_March','Bill_amt3':'Bill_amt_April','Bill_amt2':'Bill_amt_May','Bill_amt1':'Bill_amt_June'})

df2.rename(columns= {'pay_6':'Pay_January','pay_5':'Pay_February','pay_4':'Pay_March','pay_3':'Pay_April','pay_2':'Pay_May','pay_1':'Pay_June'})
df2.rename(columns= {'pay_amt6':'Pay_amt_January','pay_amt5':'Pay_amt_February','pay_amt4':'Pay_amt_March','pay_amt3':'Pay_amt_April','pay_amt2':'Pay_amt_May','pay_amt1':'Pay_amt_June'})
df2.rename(columns= {'Bill_amt6':'Bill_amt_January','Bill_amt5':'Bill_amt_February','Bill_amt4':'Bill_amt_March','Bill_amt3':'Bill_amt_April','Bill_amt2':'Bill_amt_May','Bill_amt1':'Bill_amt_June'})
```




Feature Engineering:-

Encoding Categorical Data :

- **One Hot Encoding** is a *method for converting categorical variables into a binary format*. It creates new columns for each category where **1** means the category is present and **0** means it is not. The primary purpose of One Hot Encoding is to ensure that categorical data can be effectively used in machine learning models.
- There are three columns named sex , education and marriage we applied OHE on then

```
[16] df_final = pd.get_dummies(df_final, columns=['marriage', 'education', 'sex'])  
      df2 = pd.get_dummies(df2, columns=['marriage', 'education', 'sex'])
```



Feature Engineering:-

Train Test Split

- Without adding any additional features lets check model performance

✓ Train Test Split

[+ Code](#)[+ Text](#)

```
[▶] X= df_final.drop(columns='next_month_default',axis=1)
     y= df_final['next_month_default']

     scaler = StandardScaler()
     X = scaler.fit_transform(X)

     X_train, X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

Feature Engineering:-

Feature Selection :

Feature Selection

```
[23] rf.feature_importances_
```

```
array([0.03731503, 0.04014882, 0.04086808, 0.07932968, 0.04745844,  
       0.03486227, 0.02722628, 0.02083799, 0.01582563, 0.04241196,  
       0.03739891, 0.03554222, 0.03456232, 0.03391281, 0.03449143,  
       0.0509076 , 0.04532754, 0.04139054, 0.03973833, 0.0388632 ,  
       0.04014879, 0.03938345, 0.04103 , 0.0225007 , 0.01804734,  
       0.00073806, 0.00840582, 0.01165492, 0.0010338 , 0.0053993 ,  
       0.01489638, 0.01834236])
```

```
[24] feature_scores = pd.Series(rf.feature_importances_, index=df_final.drop(columns='next_month_default',ax:  
feature_scores
```

- Selecting only important features for training model
- Scores are improved from the previous ones.

Model With Top important features

```
df_new = df_final[feature_scores.index[:20]]  
df2_new = df2[feature_scores.index[:20]]  
df_new['default'] = df_final['next_month_default']  
X = df_new.drop('default', axis=1)  
y = df_new['default']
```

```
acc = accuracy_score(y_test, test_pred)  
prec = precision_score(y_test, test_pred)  
rec = recall_score(y_test, test_pred)  
f1 = f1_score(y_test, test_pred)
```

```
print('Accuracy: ', acc)  
print('Precision: ', prec)  
print('Recall: ', rec)  
print('F1 Score: ', f1)
```

```
Accuracy: 0.8735095267363245  
Precision: 0.9025720966484801  
Recall: 0.8415697674418605  
F1 Score: 0.871004136893569
```


Feature Engineering:-

Creating new features :

✓ 1-Bill amount average

```
[33] df_final['Bill_amt_avg']=(df_final['Bill_amt_January']+df_final['Bill_amt_February']+df_final['Bill_amt_March'])  
df2['Bill_amt_avg']=(df2['Bill_amt_January']+df2['Bill_amt_February']+df2['Bill_amt_March'])
```

✓ 2-Bill Pay Value

```
[34] df_final['Bill_pay_value'] = ((df_final['Pay_amt_January'] - df_final['Bill_amt_January']) + (df_final['Pay_amt_February'] - df_final['Bill_amt_February']))  
df2['Bill_pay_value']=((df2['Pay_amt_January'] - df2['Bill_amt_January']) + (df2['Pay_amt_February'] - df2['Bill_amt_February']))
```



Model Performance

Gradient Boosting

- Gradient Boosting does not give satisfying results
- Lets try another model

```
acc = accuracy_score(y_test, test_pred)
prec = precision_score(y_test, test_pred)
rec = recall_score(y_test, test_pred)
f1 = f1_score(y_test, test_pred)
f2 = fbeta_score(y_test, test_pred, beta=2)
roc_auc = roc_auc_score(y_test, test_pred)

print('Accuracy: ', acc)
print('Precision: ', prec)
print('Recall: ', rec)
print('F1 Score: ', f1)
print('F2 Score', f2)
print('roc_auc', roc_auc)
```

```
Accuracy: 0.8183159188690842
Precision: 0.8455920709441836
Recall: 0.7853682170542635
F1 Score: 0.8143682491836222
F2 Score 0.7967167993708837
roc_auc 0.8188133822980327
```

Model Performance

CatBoost

- As we can scores are improved compared to previous model
- Lets improve it more

```
acc = accuracy_score(y_test,test_pred)
prec = precision_score(y_test,test_pred)
rec = recall_score(y_test,test_pred)
f1 = f1_score(y_test,test_pred)
f2 = fbeta_score(y_test,test_pred,beta=2)
roc_auc = roc_auc_score(y_test,test_pred)
```

```
print('Accuracy: ', acc)
print('Precision: ', prec)
print('Recall: ', rec)
print('F1 Score: ', f1)
print('F2 Score', f2)
print('roc_auc',roc_auc)
```

```
Accuracy: 0.8484326982175784
Precision: 0.8852808091562416
Recall: 0.8057170542635659
F1 Score: 0.8436271401395053
F2 Score 0.8204647491242785
roc_auc 0.8490776436778273
```

Model Performance

Applying RandomForest :

- Applying rf and evaluating performance metrics
- Results are quite good and improved
- Let try some more models

```
acc = accuracy_score(y_test,test_pred)
prec = precision_score(y_test,test_pred)
rec = recall_score(y_test,test_pred)
f1 = f1_score(y_test,test_pred)
f2 = fbeta_score(y_test,test_pred,beta=2)
roc_auc = roc_auc_score(y_test,test_pred)
```

```
print('Accuracy: ', acc)
print('Precision: ', prec)
print('Recall: ', rec)
print('F1 Score: ', f1)
print('F2 Score', f2)
print('roc_auc',roc_auc)
```

```
Accuracy: 0.8695759065765212
Precision: 0.8865641542727501
Recall: 0.8519864341085271
F1 Score: 0.8689314391599753
F2 Score 0.8586845060793984
roc_auc 0.8698414825895767
```

Model Performance

Ensemble methods - Combining random forest and GradientBoosting

- As we can see results are degraded from the previous model .
- Lets try another methods

```
▶ acc = accuracy_score(y_test,test_pred)
  prec = precision_score(y_test,test_pred)
  rec = recall_score(y_test,test_pred)
  f1 = f1_score(y_test,test_pred)
  f2 = fbeta_score(y_test,test_pred,beta=2)
  roc_auc = roc_auc_score(y_test,test_pred)

  print('Accuracy: ', acc)
  print('Precision: ', prec)
  print('Recall: ', rec)
  print('F1 Score: ', f1)
  print('F2 Score', f2)
  print('roc_auc',roc_auc)
```

```
⇒ Accuracy: 0.8544560540872772
  Precision: 0.8758937691521961
  Recall: 0.8309108527131783
  F1 Score: 0.8528095474888115
  F2 Score 0.8395339729782652
  roc_auc 0.8548115531347273
```

Model Performance

Using LGBM

- Performance metrics are good
- F2 score is also increased

```
acc = accuracy_score(y_test, test_pred)
prec = precision_score(y_test, test_pred)
rec = recall_score(y_test, test_pred)
f1 = f1_score(y_test, test_pred)
f2 = fbeta_score(y_test, test_pred, beta=2)
roc_auc = roc_auc_score(y_test, test_pred)
print('Accuracy: ', acc)
print('Precision: ', prec)
print('Recall: ', rec)
print('F1 Score: ', f1)
print('F2 Score', f2)
print('roc_auc', roc_auc)
```

```
Accuracy: 0.8867854947756607
Precision: 0.9325600215807931
Recall: 0.8374515503875969
F1 Score: 0.8824505424377792
F2 Score 0.8548889658242247
roc_auc 0.8875303671578614
```



Model Performance

Using XGBoost:

- **XGBoost** is a gradient boosting algorithm, meaning it iteratively builds a model by adding new trees that correct the errors of the previous ones.
- From this we get the highest F2 score
- Selecting this as our final model for evaluation

```
[58] f2_scorer = make_scorer(fbeta_score, beta=2)
     xgb_model = xgb.XGBClassifier()
     optimization_dict = {
         'max_depth': [2,4,6],
         'n_estimators': [50, 100, 200]
     }
     model = GridSearchCV(xgb_model, optimization_dict, scoring='accuracy', verbose=1)
     model.fit(X_train, y_train)
```

```
acc = accuracy_score(y_test, test_pred)
prec = precision_score(y_test, test_pred)
rec = recall_score(y_test, test_pred)
f1 = f1_score(y_test, test_pred)
f2 = fbeta_score(y_test, test_pred, beta=2)
roc_auc = roc_auc_score(y_test, test_pred)
```

```
print('Accuracy: ', acc)
print('Precision: ', prec)
print('Recall: ', rec)
print('F1 Score: ', f1)
print('F2 Score', f2)
print('roc_auc', roc_auc)
```

```
Accuracy:  0.8813767670559312
Precision:  0.9117417339234575
Recall:    0.8483527131782945
F1 Score:  0.8789057598192997
F2 Score 0.8603154326143566
roc_auc 0.8818753832924164
```

Model Selection

- Apply many models and comparing them on the basis of performance
- **Selecting Xgboost as out final model for output validation as it has highest f2 score**

	Accuracy	Precision	Recall	F1 Score	F2 Score	ROC AUC
Gradient Boosting	0.818316	0.845592	0.785368	0.814368	0.796717	0.818813
RandomForest	0.869576	0.886564	0.851986	0.868931	0.858685	0.869841
Ensemble (GB + RF)	0.854456	0.875894	0.830911	0.852810	0.839534	0.854812
CatBoost	0.848433	0.885281	0.805717	0.843627	0.820465	0.849078
XGBoost	0.881377	0.911742	0.848353	0.878906	0.860315	0.881875
LightGBM	0.886785	0.932560	0.837452	0.882451	0.854889	0.887530

Thank You

